

B+ Tree Recap


B Trees

- multiple keys per node, ensure each node (except root) is at least half full
- invented for disk-based storage, but still heavily used with DRAM
- B stands for Boeing, balanced, block, or Bayer (not: binary)

B+ Trees


- B+-tree: variant of B-tree where inner nodes only store pointers, no values
- leaf nodes are also often chained to simplify range scan

B+ Tree Index



- Leaf pages contain data entries, and are (optionally) chained (prev & next)
- Non-leaf pages have NO data entries

B+ Tree Basics




Parameter ***d*** = the order

Each *non-leaf* (“interior”) ***node*** has $[d, 2d]$ ***entries***

- ***Minimum 50% occupancy***

Root ***node*** has $[1, 2d]$ ***entries***


B+ Tree Basics



The n entries in a node define $n+1$ ranges

B+ Tree Basics


Non-leaf or *internal* node



For each range, in a *non-leaf* node, there is a **pointer** to another node with entries in that range

B+ Tree Basics

Non-leaf or *internal* node




Leaf nodes

Leaf nodes also have between d and $2d$ entries, and are different in that:

B+ Tree Basics

Non-leaf or *internal* node




Leaf nodes also have between d and $2d$ entries, and are different in that:

Their entry slots contain record ID to data records

B+ Tree Basics

Non-leaf or *internal* node



Leaf nodes


Leaf nodes also have between d and $2d$ entries, and are different in that:

Their entry slots contain pointers to data records

They contain a pointer to the next leaf node as well, ***for faster sequential traversal***

B+ Tree Basics

Non-leaf or *internal* node



Leaf nodes


Note that the pointers at the leaf level will be to the actual data records (rows).

We might truncate these for simpler display (as before)...


Height = 1

B+ Tree Page Format

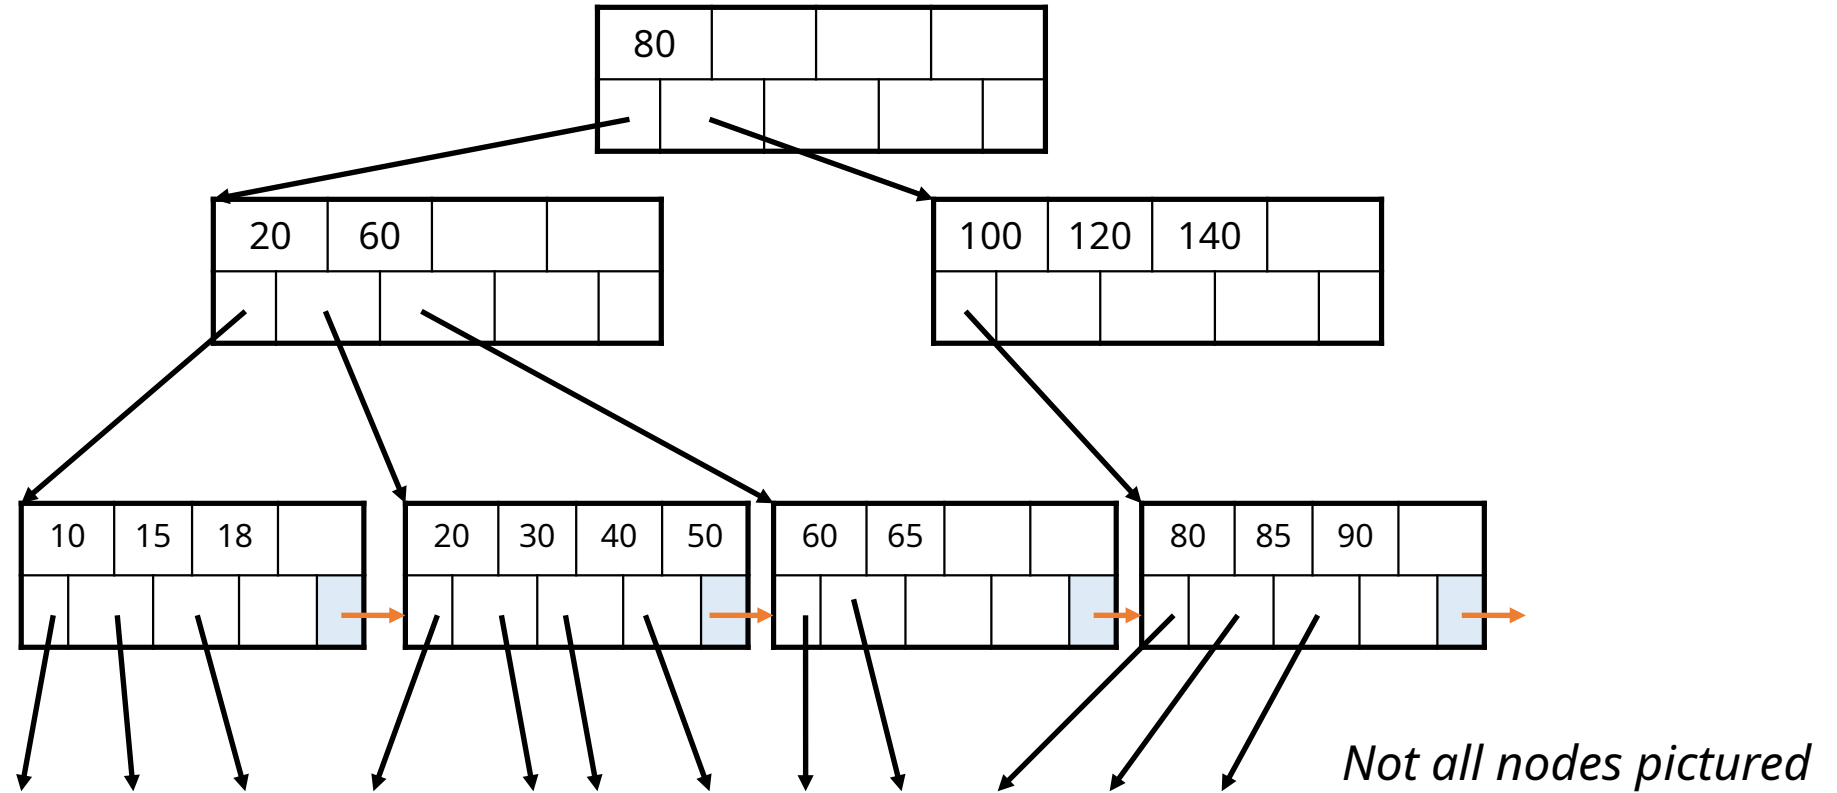
Non-leaf
Page



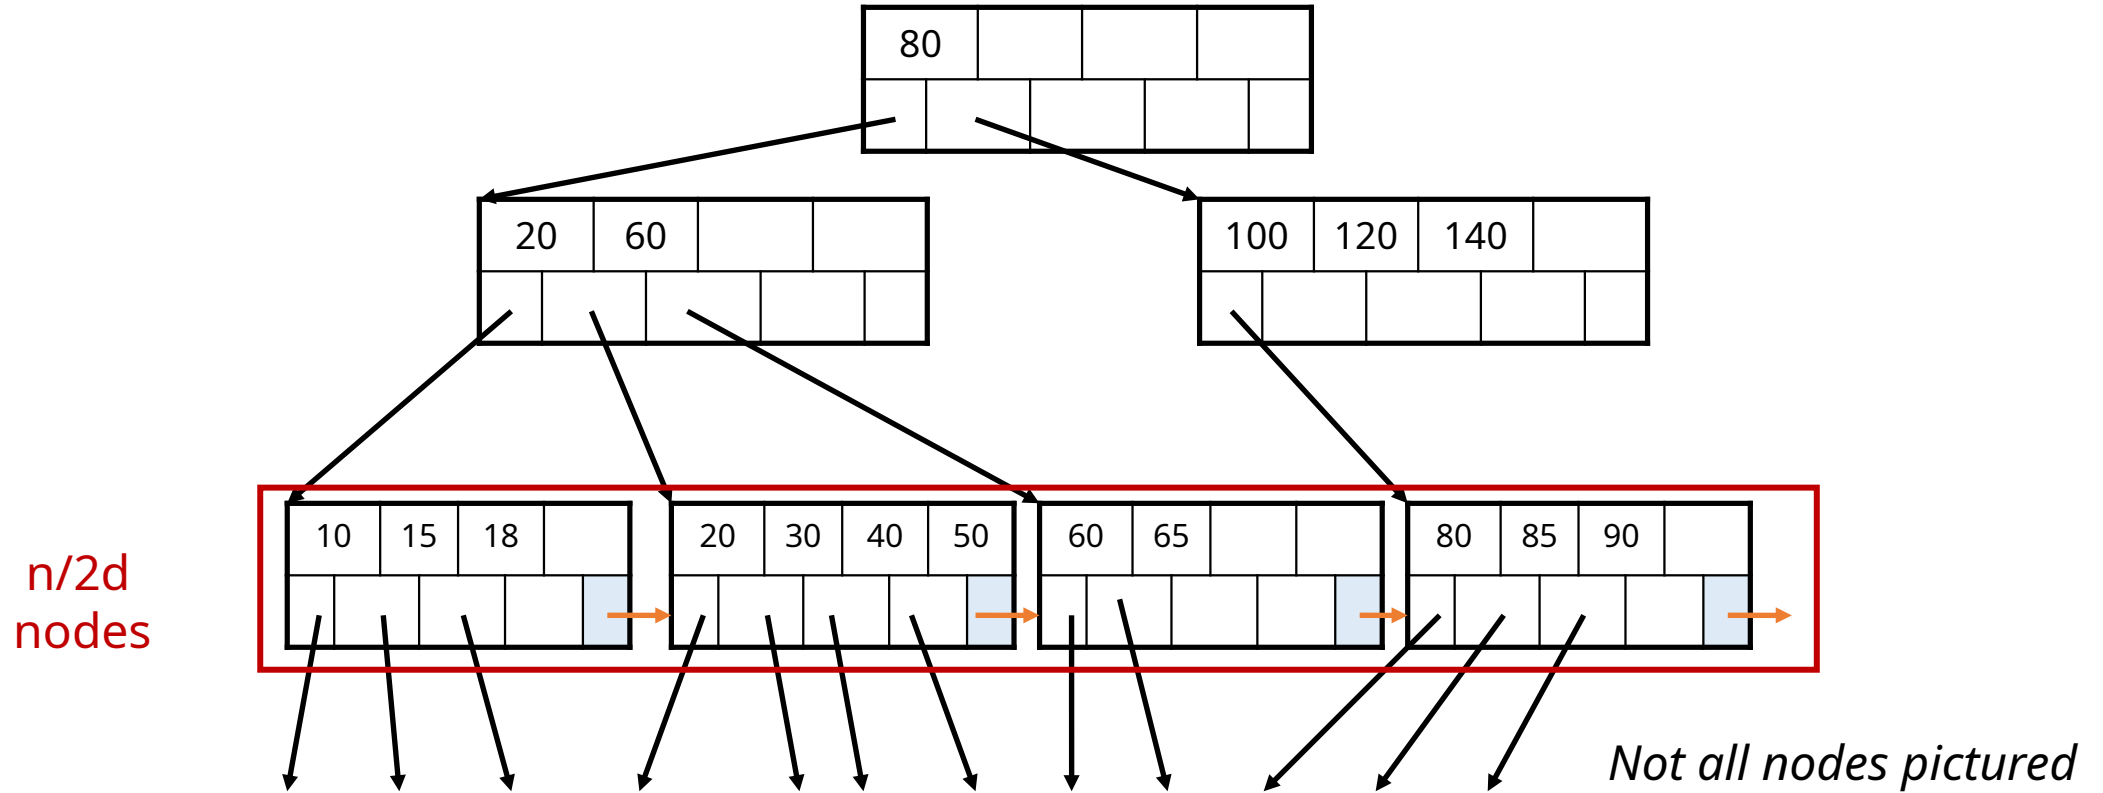
Leaf Page



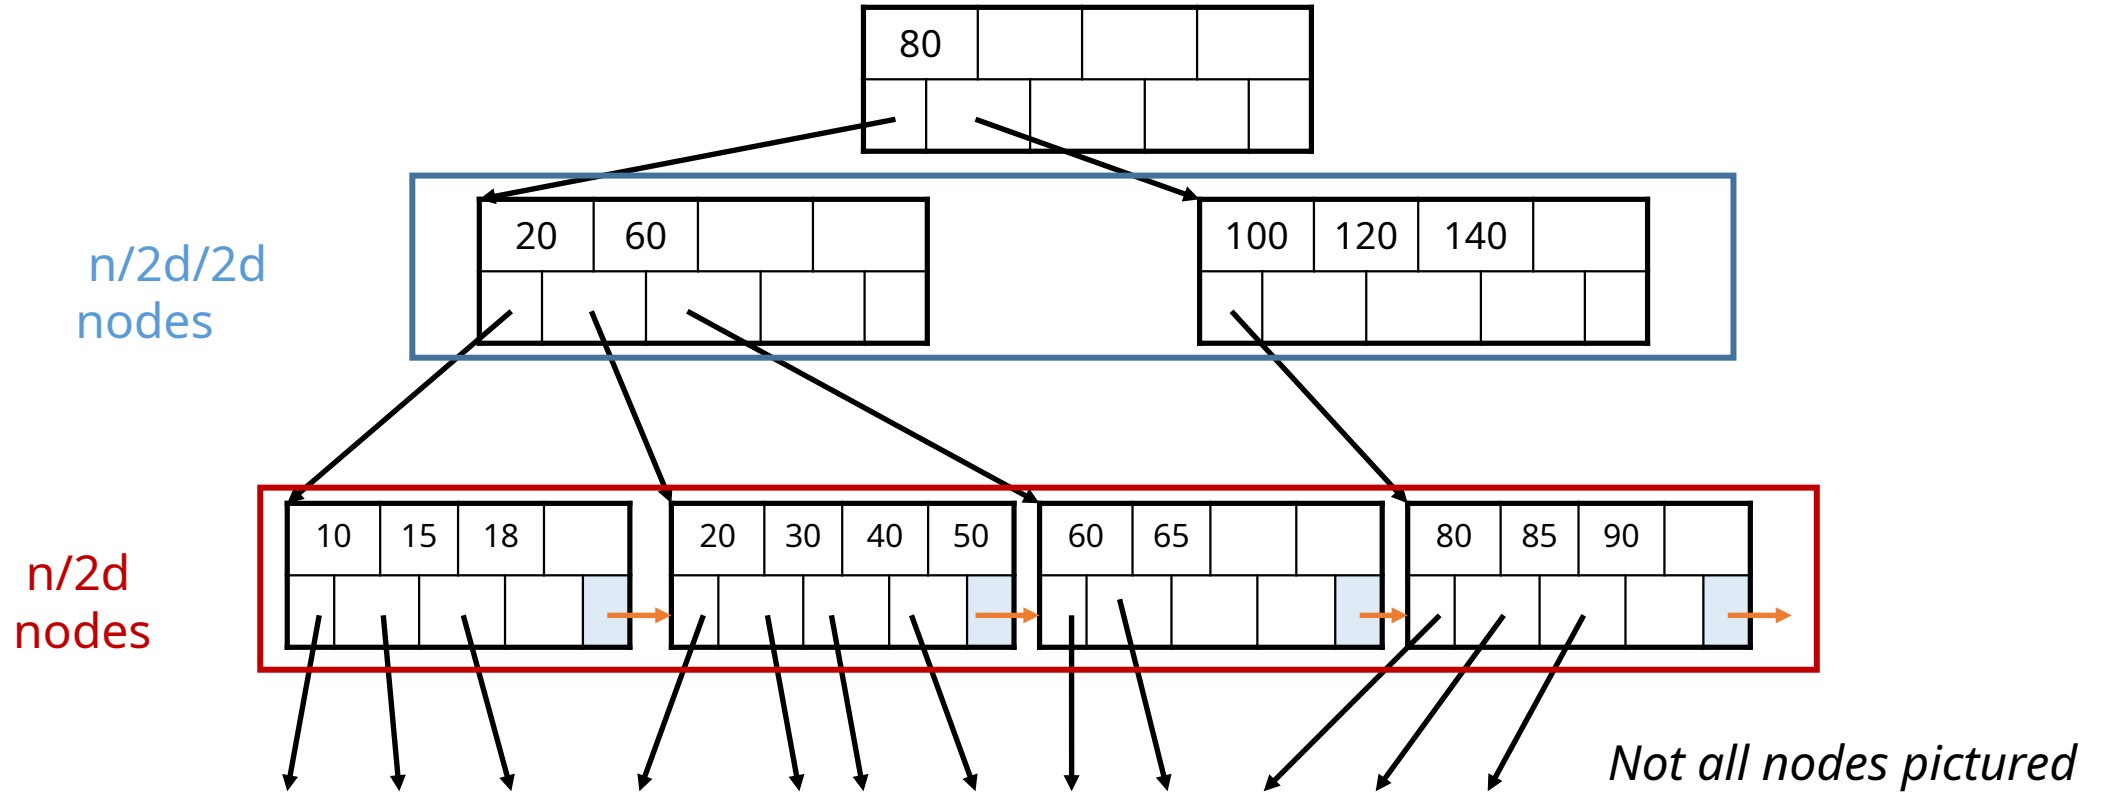
Tree height?



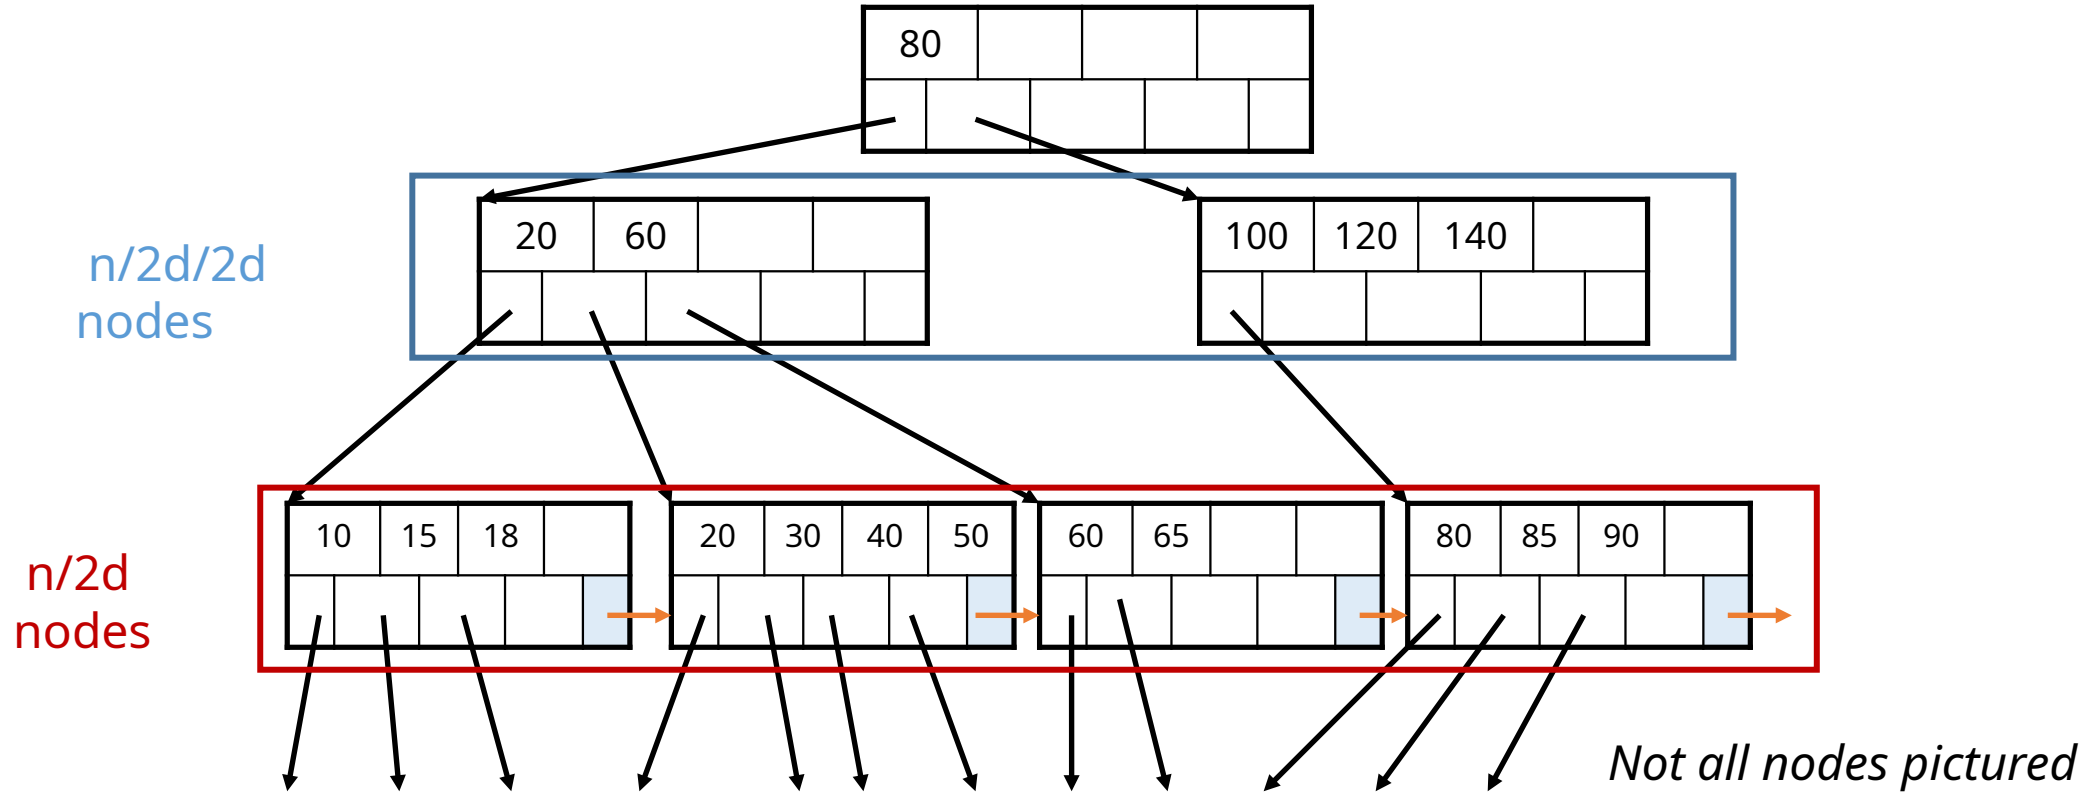
Tree height?



Tree height?



Tree height?



Best case tree height: $\log_{2d}(n)$

Tree height $\leq \log_d(n)$

B+ Trees: Operations

B+ Tree operations

A B+ tree supports the following operations:

- equality search
- range search
- insert
- delete
- bulk loading

B+ Tree: Search

- start from root
- examine index entries in non-leaf nodes to find the correct child
- traverse down the tree until a leaf node is reached

B+ Tree Exact Search Animation

K = 30?

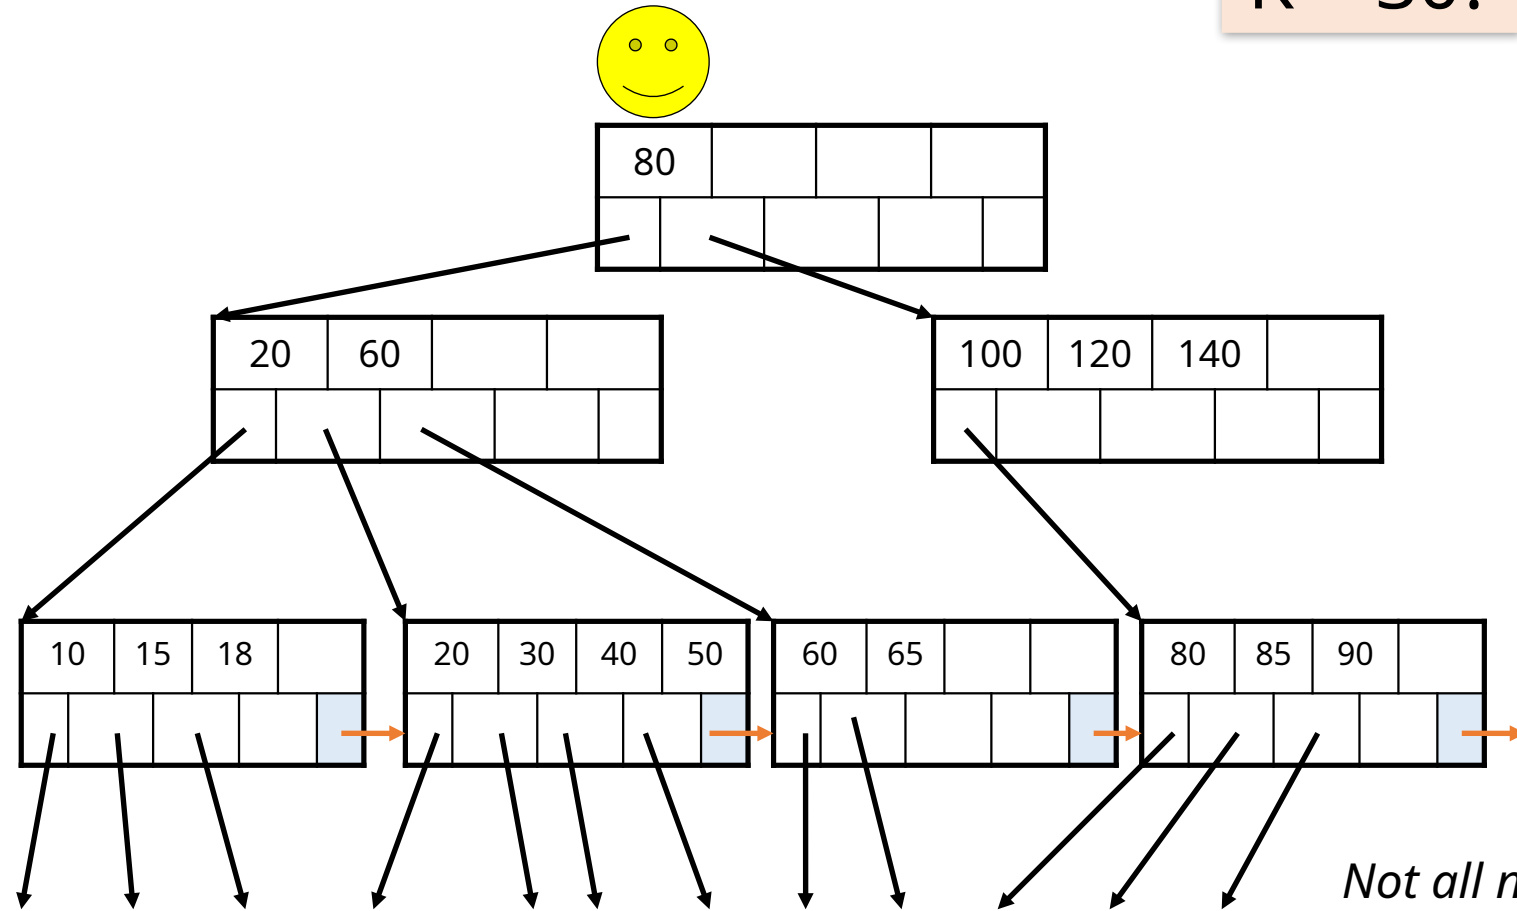
30 <

80


30 in
[20,60)

30


To the
data!



How many comparisons for the exact search?



$n/2d$ nodes

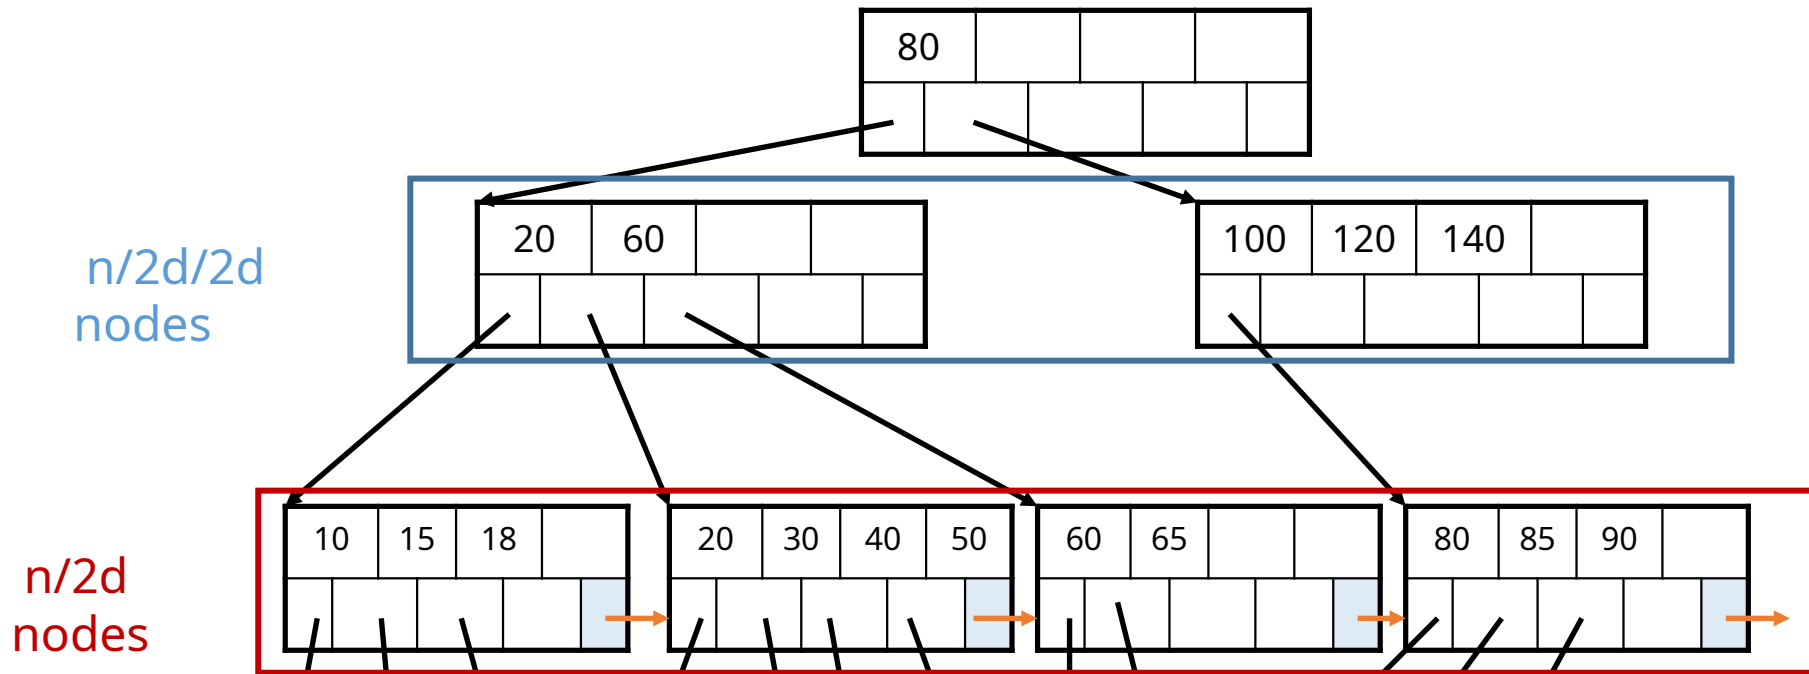


Best case tree height: $\log_{2d}(n)$
tree height $\leq \log_d(n)$

Not all nodes pictured

How many comparisons for the exact search?

comparisons: $\log_d(n) \cdot \log_2(d) = \log_2(n)$ (like binary search)



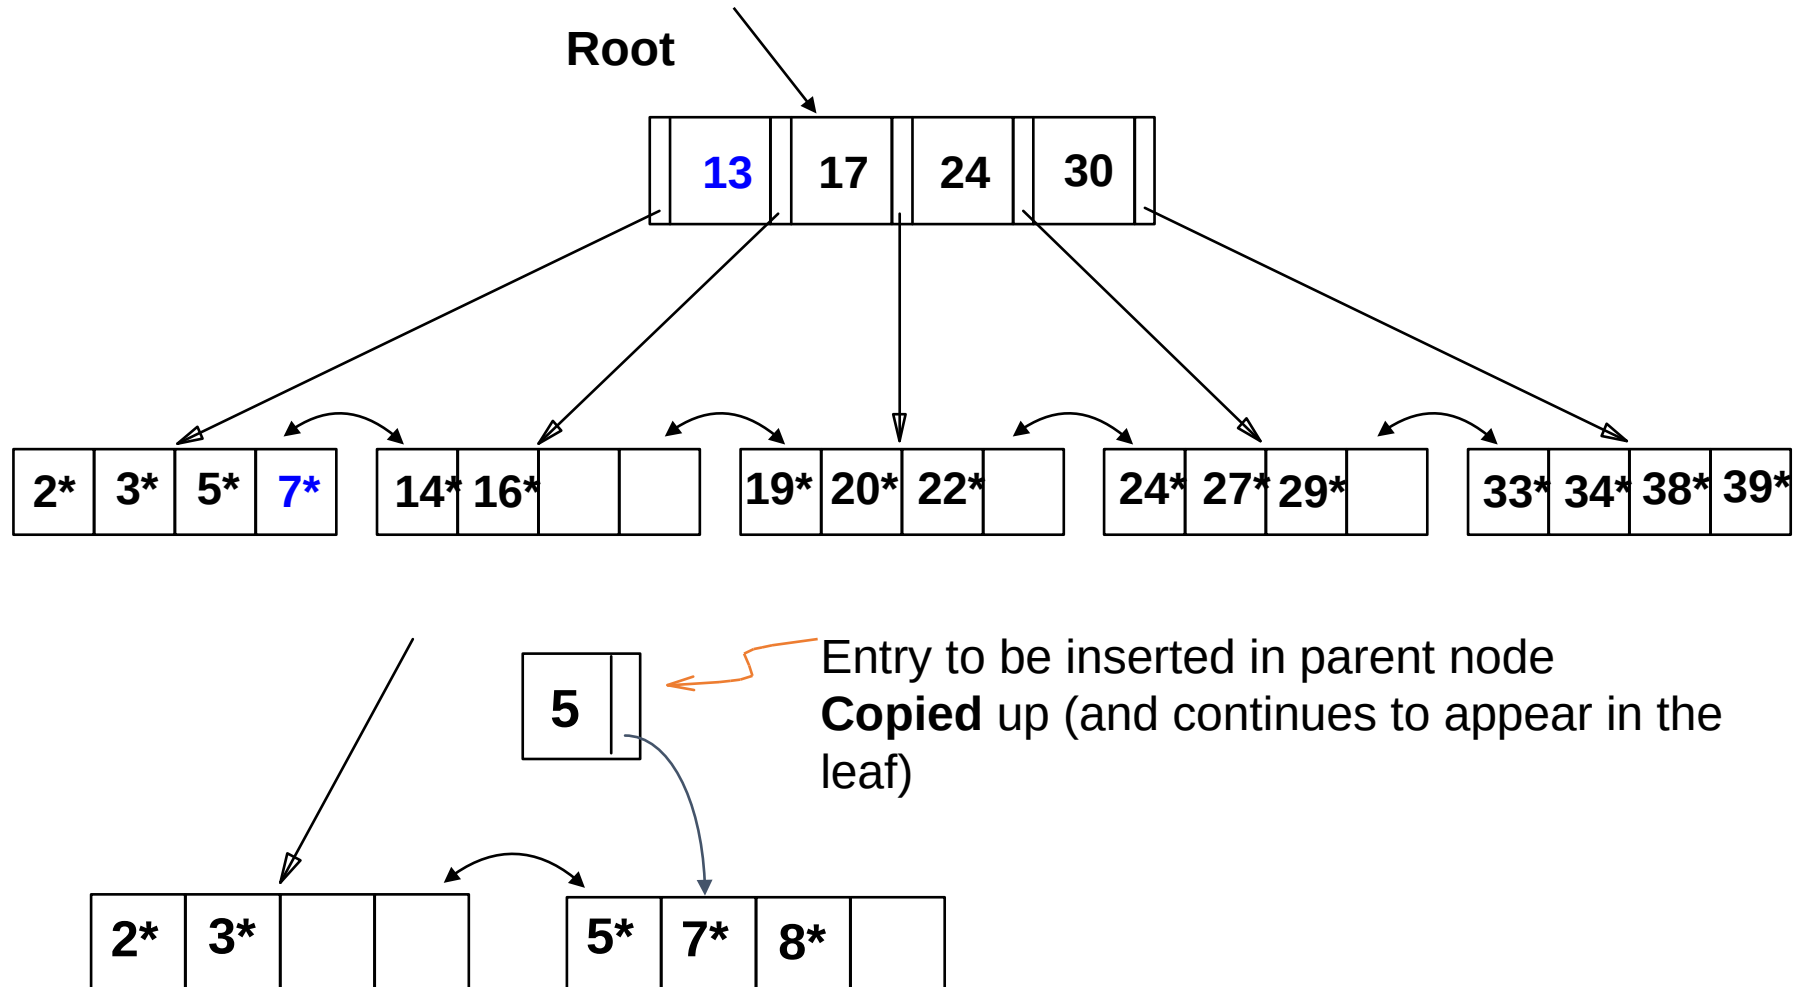
Best case tree height: $\log_{2d}(n)$

Tree height $\leq \log_d(n)$

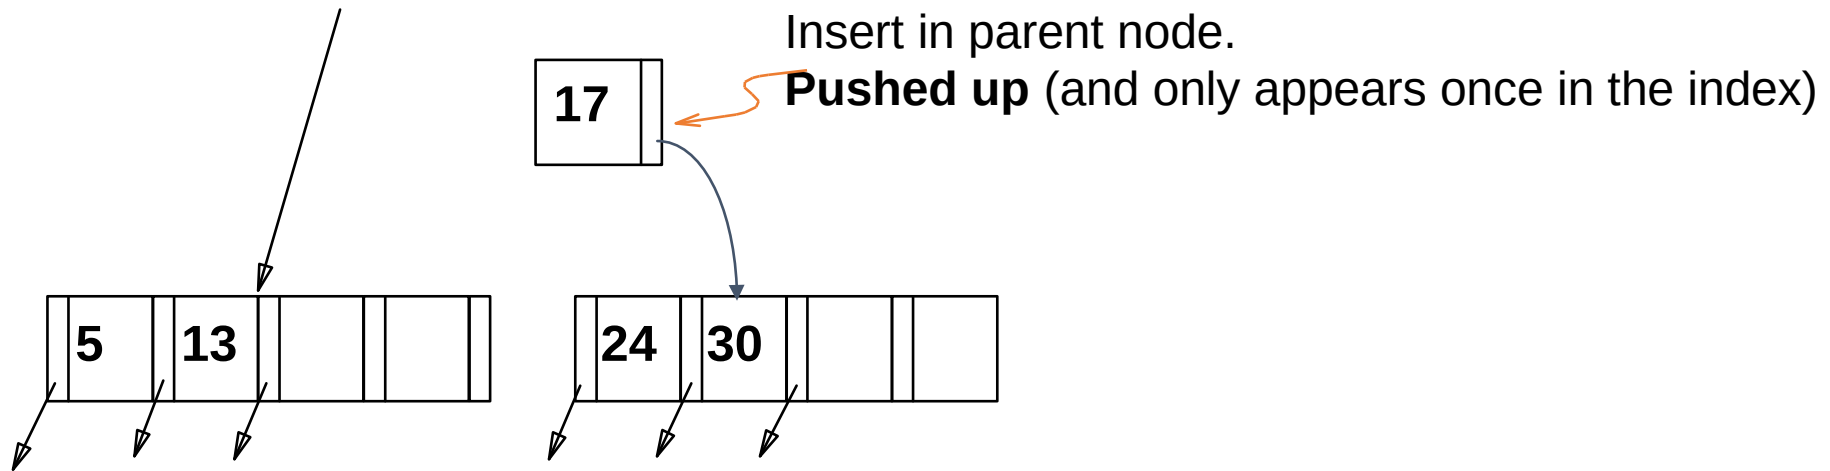
B+ Tree: Insert

- Find correct leaf L .
- Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must **split** L (into L and a new node $L2$)
 - Redistribute entries evenly, **copy up** middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- This can happen recursively
 - To split non-leaf node, redistribute entries evenly, but **pushing up** the middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
 - Tree growth: gets *wider* or *one level taller at top*.

Inserting 8* into B+ Tree




Inserting 8* into B+ Tree



**Minimum occupancy is guaranteed in both
leaf and index page splits**

Inserting 8* into B+ Tree



- Root was split: height increases by 1
- Could avoid split by re-distributing entries with a sibling
 - Sibling: immediately to left or right, and same parent


Deletion and Textbook Merge

- to avoid underfull nodes after deletion one must merge or rebalance
- B-tree delete is more complicated than insert
- if node is underfull ($<50\%$ fill factor):
 - find a neighboring node
 - if neighboring has low enough fill factor:
 - merge the two nodes into one
 - remove separator in parent
 - if parent is underfull merge parent the same way
 - else rebalance elements between the two nodes

Clustered Indexes

An index is ***clustered*** if the underlying data is ordered in the same way as the index's data entries.


Clustered vs. Unclustered Index



Clustered

Index Entries

Data Records



Unclustered

Summary

- We create **indexes** over tables in order to support ***fast (exact and range) search*** and ***insertion*** over ***multiple search keys***
- **B+ Trees** are one index data structure which support very fast exact and range search & insertion via ***high fanout***
 - ***Clustered vs. unclustered*** makes a big difference for range queries too